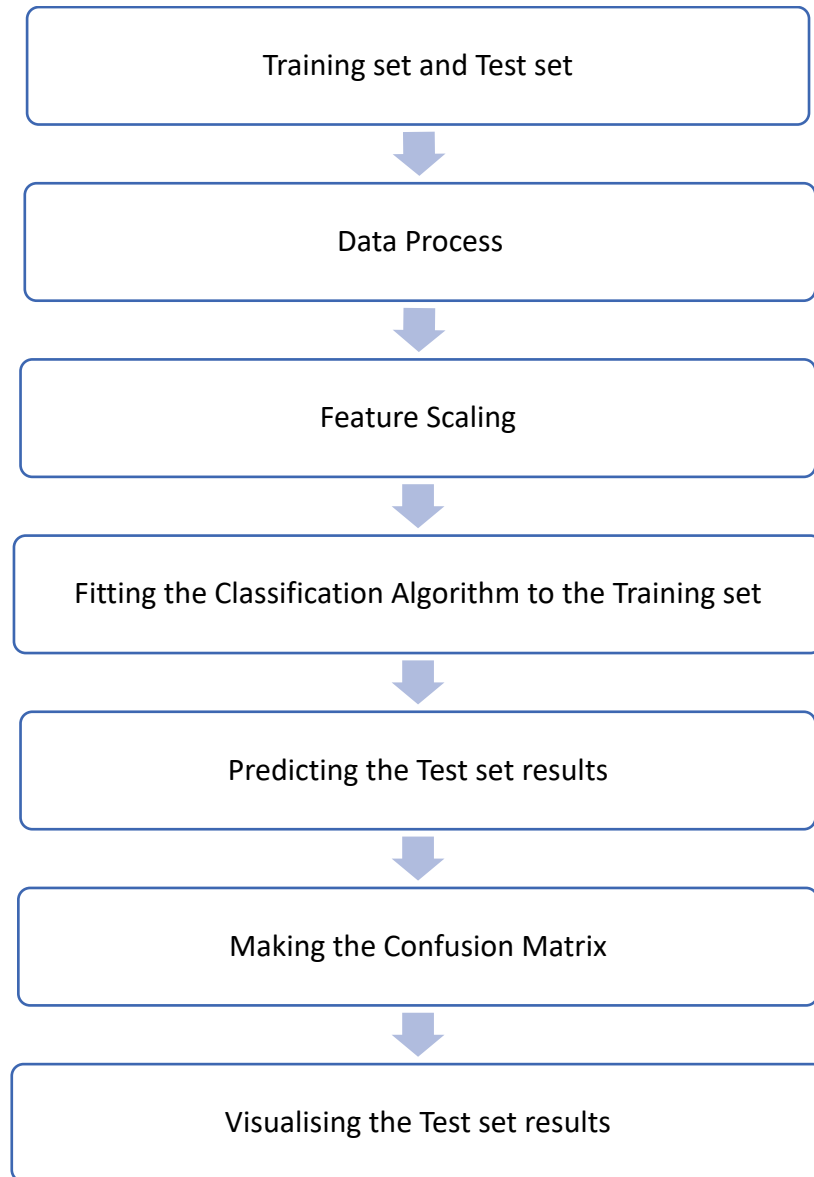


MIDTERM REPORT

**Advance in Data Sciences and Architecture
INFO 7390 - SPRING 2017**

**PROFESSOR:
SRIKANTH KRISHNAMURTHY**

**TEAM MEMBERS (Team 10):
AASHRI TANDON
PRAGATI SHAW
SARTHAK AGARWAL**



1. Programmatically downloads Q12005 and Q22005 origination data and pre-processes it.

```
##### Create Session #####
USERNAME = 'wangying0327@hotmail.com'
PASSWORD = 'Y3ZUK;\D'

payload = {
    "username": USERNAME,
    "password": PASSWORD
}

session_requests = requests.session()

login_url = "https://freddiemac.embs.com/FLoan/secure/auth.php"

result = session_requests.post(
    login_url,
    data=payload,
    headers=dict(referer=login_url)
)

url = 'https://freddiemac.embs.com/FLoan/Data/download.php'
agreement_payload = {
    "accept": "Yes",
    "action": "acceptTandC",
    "acceptSubmit": "Continue"
}
result = session_requests.post(
    url,
    agreement_payload,
    headers=dict(referer=url)
)
```

2. Builds a Logistic regression model for the
CURRENTLOANDELINQUENCYSTATUS

```

'''-----Logistic Regression-----'''
# Feature Scaling
from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)

# Fitting the Logistic Regression to the Training set
from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state = 0, solver='lbfgs')
classifier.fit(X_train, y_train)

# Predicting the Test set results
y_pred = classifier.predict(X_test)

# Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)

```

3. using Q12005 data as training data (col 4). Note anytime col 4 is > 0, add a new variable as Delinquent. Use this variable as your “Y” variable. IGNORE COL 4 AND DON’T USE IT IN YOUR MODEL.

```

def f(row):
    if row['current_loan_delinquency_status'] > 0:
        val = 1
    else:
        val = 0
    return val

# Create dummy variables
df1_dummies = pd.get_dummies(df1[['repurchase_flag', 'modification_flag']])
df2_dummies = pd.get_dummies(df2[['repurchase_flag', 'modification_flag']])

df1_d = df1.drop(['loan_sequence_no', 'repurchase_flag', 'modification_flag'], axis=1)
df2_d = df2.drop(['loan_sequence_no', 'repurchase_flag', 'modification_flag'], axis=1)

global df1_final
global df2_final
df1_final = pd.concat([df1_d, df1_dummies], axis=1)
df2_final = pd.concat([df2_d, df2_dummies], axis=1)

# create target variable
df1_final['Delinquent'] = df1_final.apply(f, axis=1)
df2_final['Delinquent'] = df2_final.apply(f, axis=1)

# Create training and testing set

X_train = df1_final.drop(['current_loan_delinquency_status', 'Delinquent'], axis=1)
y_train = df1_final['Delinquent']

X_test = df2_final.drop(['current_loan_delinquency_status', 'Delinquent'], axis=1)
y_test = df2_final['Delinquent']

X_train = preprocessing.minmax_scale(np.array(X_train).astype(float)) # scale between 0 and 1
X_test = preprocessing.minmax_scale(np.array(X_test).astype(float))

return X_train, y_train, X_test, y_test

```

4. Validates against Q22005 data and selects the best Classification model
Copyright QuantUniversity LLC. 2017 – Cannot be reused/quoted without written permission

Random Forest

```
# random Forest
def rf(X_train, y_train, X_test, y_test, return_dict_rf):
    rf = RandomForestClassifier(n_estimators=100)
    rf.fit(X_train, y_train)

    y_train_predicted = rf.predict(X_train)
    y_test_predicted = rf.predict(X_test)

    conf_mat_rf = metrics.confusion_matrix(y_test, y_test_predicted)
    print(conf_mat_rf)

    # Compute ROC curve and AUC (Area under the Curve)
    false_positive_rate, true_positive_rate, thresholds = metrics.roc_curve(y_test, y_test_predicted)

    roc_auc_rf = metrics.auc(false_positive_rate, true_positive_rate)
    return_dict_rf['roc_auc_rf'] = roc_auc_rf
    return_dict_rf['conf_mat_rf'] = conf_mat_rf

    # Plot ROC Curve
    plt.title("Random Forest")
    plt.plot(false_positive_rate, true_positive_rate, 'b',
             label='AUC = %0.2f' % roc_auc_rf)
    plt.legend(loc='lower right')
    plt.plot([0, 1], [0, 1], 'r--')
    plt.xlim([-0.1, 1.2])
    plt.ylim([-0.1, 1.2])
    plt.ylabel('True Positive Rate')
    plt.xlabel('False Positive Rate')
    plt.show()
```

Neural Network

```

def nn(X_train, y_train, X_test, y_test, return_dict_nn):

    nn = MLPClassifier()

    nn.fit(X_train, y_train)

    y_train_predicted = nn.predict(X_train)
    y_test_predicted = nn.predict(X_test)

    conf_mat_nn = metrics.confusion_matrix(y_test, y_test_predicted)
    print(conf_mat_nn)

    # Compute ROC curve and AUC (Area under the Curve)
    false_positive_rate, true_positive_rate, thresholds = metrics.roc_curve(y_test, y_test_predicted)

    roc_auc_nn = metrics.auc(false_positive_rate, true_positive_rate)
    return_dict_nn['roc_auc_nn'] = roc_auc_nn
    return_dict_nn['conf_mat_nn'] = conf_mat_nn

    # Plot ROC Curve
    plt.title("Neural Network")
    plt.plot(false_positive_rate, true_positive_rate, 'b',
             label='AUC = %0.2f' % roc_auc_nn)
    plt.legend(loc='lower right')
    plt.plot([0, 1], [0, 1], 'r--')
    plt.xlim([-0.1, 1.2])
    plt.ylim([-0.1, 1.2])
    plt.ylabel('True Positive Rate')
    plt.xlabel('False Positive Rate')
    plt.show()

```

5. H2O

```
: import h2o
```

```
h2o.init(strict_version_check=False,port=9957)
```

Checking whether there is an H2O instance running at <http://localhost:9957>.... not found.

Attempting to start a local H2O server...

Java Version: openjdk version "1.8.0_121"; OpenJDK Runtime Environment (Zulu 8.20.0.5-macosx) (build 1.8.0_121-b15); OpenJDK 64-Bit Server VM (Zulu 8.20.0.5-macosx) (build 25.121-b15, mixed mode)

Starting server from /anaconda3/lib/python3.6/site-packages/h2o/backend/bin/h2o.jar

Ice root: /var/folders/hf/c05h_6kd5wlds2ysvlyvpylc0000gn/T/tmpdo3zi7w5

JVM stdout: /var/folders/hf/c05h_6kd5wlds2ysvlyvpylc0000gn/T/tmpdo3zi7w5/h2o_wangying_started_from_python.out

JVM stderr: /var/folders/hf/c05h_6kd5wlds2ysvlyvpylc0000gn/T/tmpdo3zi7w5/h2o_wangying_started_from_python.err

Server is running at <http://127.0.0.1:9957>

Connecting to H2O server at <http://127.0.0.1:9957>... successful.

H2O cluster uptime:	01 secs
H2O cluster timezone:	America/New_York
H2O data parsing timezone:	UTC
H2O cluster version:	3.22.0.2
H2O cluster version age:	15 days
H2O cluster name:	H2O_from_python_wangying_kbgw7l
H2O cluster total nodes:	1
H2O cluster free memory:	3.556 Gb
H2O cluster total cores:	8
H2O cluster allowed cores:	8
H2O cluster status:	accepting new members, healthy
H2O connection url:	http://127.0.0.1:9957
H2O connection proxy:	None
H2O internal security:	False
H2O API Extensions:	XGBoost, Algos, AutoML, Core V3, Core V4
Python version:	3.6.6 final

Out[27]:

	model_id	auc	logloss	mean_per_class_error	rmse	mse
0	StackedEnsemble_AllModels_AutoML_20181207_133501	0.911731	0.139201	0.210976	0.187376	0.035110
1	StackedEnsemble_BestOfFamily_AutoML_20181207_1...	0.911731	0.139201	0.210976	0.187376	0.035110
2	DRF_1_AutoML_20181207_133501	0.907181	0.155591	0.219043	0.203407	0.041374
3	XRT_1_AutoML_20181207_133501	0.906353	0.157899	0.215183	0.205562	0.042256
4	GLM_grid_1_AutoML_20181207_133501_model_1	0.745118	0.221634	0.329522	0.241570	0.058356

```
In [28]: model_set=aml_leaderboard_df['model_id']  
mod_best=h2o.get_model(model_set[0])
```

Best Model:

StackedEnsemble_AllModels_AutoML_20181207_133501

6. Parameterize the input (example it should take Q12005) and modify the code so that it outputs the 5 parameters listed in the matrix below.

matrix_classification

	Quarter	No_of_actual_delq	No_of_pred_delq	No_of_records	No_of_delq_properly_classified	No_of_nonDelq_improperly_classified_as_delq
0	Q12005	8500	3761	200000	741	3020
0	Q12005	8500	3761	200000	741	3020

7. Write another script that calls the above classification script from Q11999-Q42016 and computes the following matrix.

```
arg_len = len(sys.argv)
quarters = []
end = ''
if arg_len == 3:
    startQuarter = sys.argv[1]
    endQuarter = sys.argv[2]
    end = endQuarter
    while(startQuarter != endQuarter):
        print(startQuarter)
        quarters.append(startQuarter)
        startQuarter = get_next_quarter(startQuarter)
        quarters.append(endQuarter)
elif arg_len == 2:
    startQuarter = sys.argv[1]
    print("----" + startQuarter)
    quarters.append(startQuarter)
    end = startQuarter
else:
    print("running for default Q12005")
    quarter = 'Q12005'
    quarters.append(quarter)
    end = quarter
```

8. • Extra credit: Note that each invocation of is independent. There is scope to parallelize it. Parallelize this so that you can run two or more models concurrently (task parallel). Review R and Python packages that can be used to parallelize it to generate the below matrix as a dataframe and export it to a csvfile.


```

for q in quarters:
    nextQuarter = get_next_quarter(q)
    load_data_df(q, nextQuarter)
    remove_nan(df1)
    remove_nan(df2)
    X_train, y_train, X_test, y_test = process_data()
    print("process data executed")

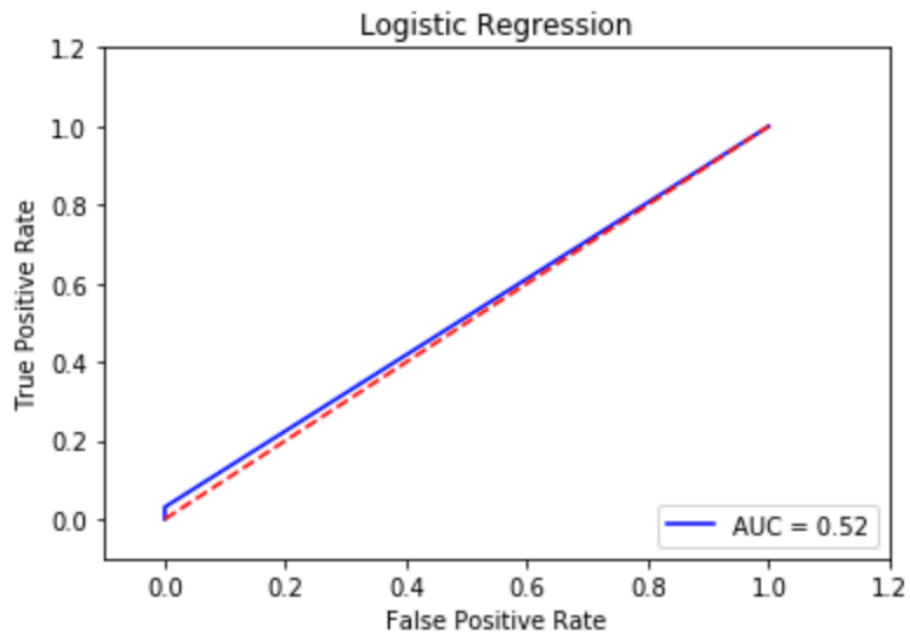
    p1 = Process(target=logred,args=(X_train, y_train,X_test,y_test,return_dict_logred))
    p2 = Process(target=rf,args=(X_train, y_train,X_test,y_test,return_dict_rf))
    p3 = Process(target=nn,args=(X_train, y_train,X_test,y_test,return_dict_nn))
    p1.start()
    p2.start()
    p3.start()
    p1.join()
    p2.join()
    p3.join()
    matrix=genMatrix(q,return_dict_logred['roc_auc_logred'],return_dict_rf['roc_auc_rf'],return_dict_nn['
        roc_auc_nn'],return_dict_logred['conf_mat_logred'],return_dict_rf['conf_mat_rf'],return_dict_nn['
        conf_mat_nn'],matrix)

print(matrix)
matrix.to_csv('matrix_classification.csv')

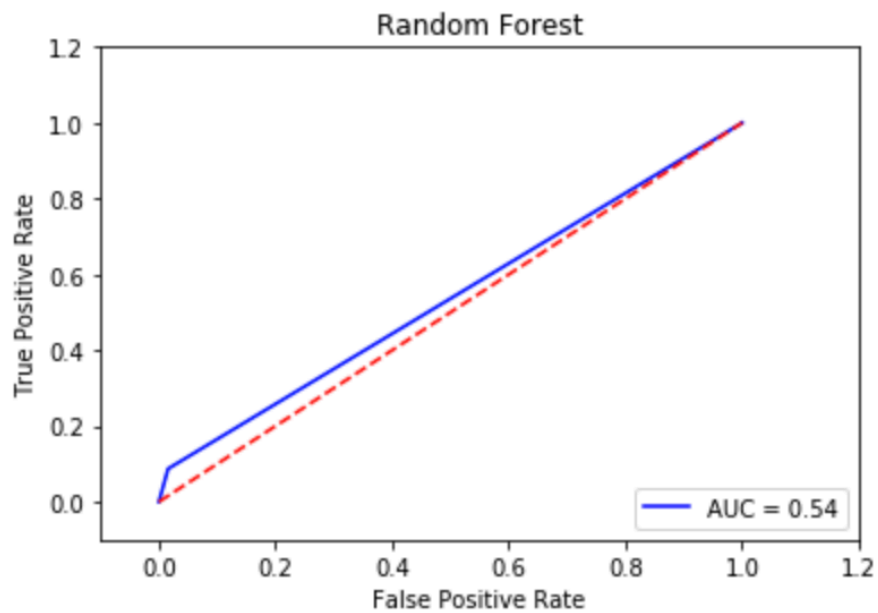
```

Graph :


```
[[191498    2]
 [  8239   261]]
```



```
[[188480    3020]
 [  7759    741]]
```



```
[[191498      2]
 [  8239    261]]
```



	Quarter	No_of_actual_delq	No_of_pred_delq	No_of_records	\
0	Q12005	8500.0	3761.0	200000.0	
0	Q12005	8500.0	3761.0	200000.0	
		No_of_delq_properly_classified	No_of_nonDelq_improperly_classified_as_delq		
0		741.0		3020.0	
0		741.0		3020.0	

```
]:
```